

EXP NO: 03
DATE: 30/08/2025

SGD WITH MOMENTUM VS ADAM OPTIMIZER

AIM: To implement a training algorithm using Stochastic Gradient Descent (SGD) with momentum and compare it with the Adam optimizer using the CIFAR-10 dataset by analysing their convergence rates and classification performance.

ALGORITHM:

- Load CIFAR-10 dataset and preprocess images with normalization.
- Define a Simple CNN model for image classification.
- Train the model twice: Using SGD with Momentum and Adam Optimizer
- Perform forward propagation to compute predictions.
- Calculate loss using Cross-Entropy Loss.
- Update weights using selected optimizer.
- Compare both models using: Training Loss Curve, Test Accuracy Curve.

CODE:

```
import torch
import torch.nn as nn
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms
import matplotlib.pyplot as plt

# Data Loading & Normalization
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])
```

```
trainset = torchvision.datasets.CIFAR10(root='./data', train=True, download=True,  
transform=transform)
```

```
trainloader = torch.utils.data.DataLoader(trainset, batch_size=128, shuffle=True,  
num_workers=2)
```

```
testset = torchvision.datasets.CIFAR10(root='./data', train=False, download=True,  
transform=transform)
```

```
testloader = torch.utils.data.DataLoader(testset, batch_size=128, shuffle=False,  
num_workers=2)
```

```
class SimpleCNN(nn.Module):
```

```
    def __init__(self):  
        super(SimpleCNN, self).__init__()  
        self.conv1 = nn.Conv2d(3, 32, 3, padding=1)  
        self.conv2 = nn.Conv2d(32, 64, 3, padding=1)  
        self.pool = nn.MaxPool2d(2, 2)  
        self.fc1 = nn.Linear(64 * 8 * 8, 128)  
        self.fc2 = nn.Linear(128, 10)  
        self.relu = nn.ReLU()
```

```
    def forward(self, x):
```

```
        x = self.pool(self.relu(self.conv1(x)))  
        x = self.pool(self.relu(self.conv2(x)))  
        x = x.view(-1, 64 * 8 * 8)  
        x = self.relu(self.fc1(x))  
        x = self.fc2(x)  
        return x
```

```
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

def train_model(optimizer_type="sgd", epochs=10, lr=0.01, momentum=0.9):
    model = SimpleCNN().to(device)
    criterion = nn.CrossEntropyLoss()

    if optimizer_type == "sgd":
        optimizer = optim.SGD(model.parameters(), lr=lr, momentum=momentum)
    elif optimizer_type == "adam":
        optimizer = optim.Adam(model.parameters(), lr=lr)

    train_losses = []
    test_accuracies = []

    for epoch in range(epochs):
        model.train()
        running_loss = 0.0
        for inputs, labels in trainloader:
            inputs, labels = inputs.to(device), labels.to(device)
            optimizer.zero_grad()
            outputs = model(inputs)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()
            running_loss += loss.item()

        train_losses.append(running_loss / len(trainloader))
        test_accuracies.append(test_accuracy(model, testloader))

    return train_losses, test_accuracies
```

```
train_losses.append(running_loss / len(trainloader))

model.eval()

correct, total = 0, 0

with torch.no_grad():

    for inputs, labels in testloader:

        inputs, labels = inputs.to(device), labels.to(device)

        outputs = model(inputs)

        _, predicted = torch.max(outputs, 1)

        total += labels.size(0)

        correct += (predicted == labels).sum().item()

    test_accuracies.append(100 * correct / total)

    print(f"Epoch {epoch+1}/{epochs} | Loss: {train_losses[-1]:.4f} | Accuracy: {test_accuracies[-1]:.2f}%")

return train_losses, test_accuracies

# Train with both optimizers

sgd_losses, sgd_acc = train_model(optimizer_type="sgd", epochs=10, lr=0.01,
momentum=0.9)

adam_losses, adam_acc = train_model(optimizer_type="adam", epochs=10, lr=0.001)

# Plot training loss comparison

plt.figure(figsize=(10,5))

plt.plot(sgd_losses, label='SGD + Momentum')
```

```
plt.plot(adam_losses, label='Adam')

plt.xlabel('Epochs'); plt.ylabel('Training Loss')

plt.title('Training Loss Comparison')

plt.legend()

plt.show()
```

```
# Plot accuracy comparison

plt.figure(figsize=(10,5))

plt.plot(sgd_acc, label='SGD + Momentum')

plt.plot(adam_acc, label='Adam')


plt.xlabel('Epochs'); plt.ylabel('Test Accuracy (%)')


plt.title('Test Accuracy Comparison')

plt.legend()

plt.show()
```

OUTPUT:





RESULT: The Adam optimizer achieved faster convergence and higher accuracy 72.78% compared to SGD with Momentum 72.08% on CIFAR-10.