

EXP NO: 04
DATE: 06/09/2025

IMPLEMENT CNN FROM SCRATCH

AIM: To implement a Convolutional Neural Network (CNN) from scratch for image classification using the CIFAR-10 dataset, evaluate its performance, and visualize the learned convolution filters.

ALGORITHM:

- Load the CIFAR-10 dataset and apply normalization preprocessing.
- Define a CNN model with convolution, ReLU, pooling, and fully-connected layers.
- Perform forward propagation to generate class predictions.
- Compute the classification loss using Cross-Entropy Loss.
- Perform backpropagation and update model weights using Adam optimizer.
- Evaluate classification accuracy on the test dataset.
- Visualize the learned filters from the first convolution layer.

CODE:

```
# CNN from scratch on CIFAR-10

import torch, torch.nn as nn, torch.optim as optim

import torch.nn.functional as F

from torchvision import datasets, transforms

from torch.utils.data import DataLoader

import matplotlib.pyplot as plt

import numpy as np

import torchvision

# Device

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

# Data Preprocessing

transform = transforms.Compose([
```

```
transforms.ToTensor(),  
  
transforms.Normalize((0.5,0.5,0.5),(0.5,0.5,0.5))  
  
])  
  
trainset = datasets.CIFAR10(root='./data', train=True, download=True,  
transform=transform)  
  
testset = datasets.CIFAR10(root='./data', train=False, download=True,  
transform=transform)  
  
trainloader = DataLoader(trainset, batch_size=64, shuffle=True)  
testloader = DataLoader(testset, batch_size=64, shuffle=False)  
  
  
# CNN Model  
  
class SimpleCNN(nn.Module):  
  
    def __init__(self):  
        super(SimpleCNN, self).__init__()  
  
        self.conv1 = nn.Conv2d(3, 32, 3, padding=1)  
        self.conv2 = nn.Conv2d(32, 64, 3, padding=1)  
  
        self.pool = nn.MaxPool2d(2,2)  
        self.fc1 = nn.Linear(64*8*8, 128)  
        self.fc2 = nn.Linear(128, 10)  
  
  
    def forward(self, x):  
        x = self.pool(F.relu(self.conv1(x))) # 32x16x16  
        x = self.pool(F.relu(self.conv2(x))) # 64x8x8  
        x = x.view(-1, 64*8*8)  
        x = F.relu(self.fc1(x))  
        return self.fc2(x)  
  
  
model = SimpleCNN().to(device)
```

```
criterion = nn.CrossEntropyLoss()  
optimizer = optim.Adam(model.parameters(), lr=0.001)
```

```
# Training  
for epoch in range(10):  
    model.train()  
    for imgs, labels in trainloader:  
        imgs, labels = imgs.to(device), labels.to(device)  
        optimizer.zero_grad()  
        outputs = model(imgs)  
        loss = criterion(outputs, labels)  
        loss.backward()  
        optimizer.step()  
    print(f"Epoch {epoch+1} complete")
```

```
# Evaluation  
correct, total = 0, 0  
model.eval()  
with torch.no_grad():  
    for imgs, labels in testloader:  
        imgs, labels = imgs.to(device), labels.to(device)  
        outputs = model(imgs)  
        _, pred = torch.max(outputs, 1)  
        total += labels.size(0)  
        correct += (pred == labels).sum().item()
```

```
accuracy = 100 * correct / total  
print(f"Test Accuracy: {accuracy:.2f}%")
```

```
# Filter Visualization

weights = model.conv1.weight.data.cpu()

grid = torchvision.utils.make_grid(weights, nrow=8, normalize=True, pad_value=1)

plt.figure(figsize=(8,8))


plt.imshow(np.transpose(grid.numpy(), (1,2,0)))

plt.title("Learned Filters in Conv1")

plt.axis("off")

plt.show()
```

OUTPUT:



RESULT: The CNN model achieved a test accuracy of 72% on the CIFAR-10 dataset, and the visualization clearly shows the learned edge- and texture-detecting filters from the first convolution layer.